

Schedule for First International Workshop on Composable Data Management Systems, 9th September 2022 @ VLDB

8:30 AM to 8:45 AM	Welcome and Kickoff - Mohamed & Satya
8:45 AM to 9:30 AM	Keynote 1 Title: <i>From parsing, to query processing, to resiliency: Deriving unexpected value from database technology through componentization</i> Speaker: <i>Jonathan Goldstein, Researcher, Microsoft</i>
9:30 AM to 10:00 AM	Keynote 2 Title: <i>Evolving an Organization's Data Management Discipline in support of their Machine Learning Activities</i> Speaker: <i>David Cohen, Senior Principal Engineer, Intel</i>
10:00 AM to 10:25 AM	Break
10:25 AM to 11:30 AM	Invited Talks 1-3 (20 mins each) - 5 mins overflow
	Title: <i>SAP's global data mesh for composable data management supporting cross-domain consumption</i> Speakers: <i>Anil K Goel, CTO / Head of Technology Office, SAP HANA Database and Analytics & Mihnea Andrei, SAP HANA Database and Analytics</i>
	Title: <i>GoogleSQL: A SQL Language as a Component</i> Speaker: <i>Jeff Shute, Distinguished Engineer, Google</i>
	Title: <i>The Single Product Approach: Supporting Diverse Workloads with Snowflake's Query Engine</i> Speaker: <i>Jiaqi Yan, Principal Software Engineer, Snowflake Computing</i>
11:30 AM to 12:15 PM	Keynote 3 Title: <i>Substrait: Rethinking DBMS Composability</i> Speaker: <i>Jacques Nadeau, Co-creator of Apache Arrow, Dremio and Substrait</i>
12:15 PM to 1:15 PM	Lunch

1:15 PM to 2:20 PM	Invited Talks 4-6 (20 mins each) - 5 mins overflow
	Title: <i>Coral: A SQL translation and rewrite engine for modern data lakes</i> Speaker: <i>Walaa Eldin Moustafa, Senior Staff Software Engineer, LinkedIn</i>
	Title: <i>Delta Lake: The one storage format for all your analytics</i> Speaker: <i>Tathagata Das, Staff Software Engineer, Databricks, Apache Spark PMC</i>
	Title: <i>Velox: Unifying execution engines across Meta and beyond</i> Speaker: <i>Pedro Pedreira, Software Engineer, Meta Platforms Inc.</i>
2:20 PM to 2:55 PM	Short Paper Presentations 1-3 (10 mins each) - 5 mins overflow
	Title: <i>Assembling a Query Engine From Spare Parts</i> Speakers: <i>Mosha Pasumansky, Firebolt Analytics & Benjamin Wagner, Firebolt Analytics</i>
	Title: <i>Hamilton: a modular open source declarative paradigm for high level modeling of dataflows</i> Speaker: <i>Stefan Krawczyk, Stitch Fix</i>
	Title: <i>Efficient Property Projections of Graph Queries over Relational Data</i> Speaker: <i>Danica Porobic, Oracle</i>
2:55 PM to 3:15 PM	Break
3:15 PM to 4:00 PM	Keynote 4 Title: <i>Shared, High-Performance Software Components for Shared, High-Performance Hardware</i> Speaker: <i>Xiaosong Ma, Principal scientist, Qatar Computing Research Institute, HBKU</i>
4:00 PM to 4:45 PM	Keynote 5 Title: <i>DuckDB - A Modern Modular & Extensible Database System</i> Speaker: <i>Mark Raasveldt, CTO, DuckDB Labs & Postdoctoral Researcher, CWI</i>

4:45 PM to 5:20 PM	Short Paper Presentations 4-6 (10 mins each) - 5 mins overflow
	Title: <i>Unifying Metadata for Stream and Batch Queries in a Cloud Service</i> Speaker: <i>Daniel Bauer, IBM Research Europe</i>
	Title: <i>Towards a Modular Data Management System Framework</i> Speaker: <i>Haralampos Gavriilidis, Technische Universität Berlin</i>
	Title: <i>Knowledge-based Recommendation for Polyglot Persistence</i> Speaker: <i>Philipp Eisenhuth, University of Bayreuth</i>
5:20 PM to 5:30 PM	Closing Remarks - Mohamed & Satya