

# Schedule for First International Workshop on Composable Data Management Systems, 9th September 2022 @ VLDB

<b>8:30 - 8:40</b>	<b>Welcome and Kickoff - Mohamed &amp; Satya</b>
<b>8:40 - 9:20</b>	<b>Keynote 1</b>  <b>Title:</b> <i>From parsing, to query processing, to resiliency: Deriving unexpected value from database technology through componentization</i>  <b>Speaker:</b> <i>Jonathan Goldstein, Researcher, Microsoft</i>
<b>9:20 - 10:00</b>	<b>Keynote 2</b>  <b>Title:</b> <i>Evolving an Organization's Data Management Discipline in support of their Machine Learning Activities</i>  <b>Speaker:</b> <i>David Cohen, Senior Principal Engineer, Intel</i>
<b>10:00 - 10:20</b>	<b>Invited Talk 1</b>  <b>Title:</b> <i>SAP's global data mesh for composable data management supporting cross-domain consumption</i>  <b>Speakers:</b> <i>Anil K Goel, CTO / Head of Technology Office, SAP HANA Database and Analytics &amp; Mihnea Andrei, SAP HANA Database and Analytics</i>
<b>10:20 - 10:30</b>	<b>Short Paper 1</b>  <b>Title:</b> <i>Assembling a Query Engine From Spare Parts</i>  <b>Speakers:</b> <i>Mosha Pasumansky, Firebolt Analytics &amp; Benjamin Wagner, Firebolt Analytics</i>
<b>10:30 - 11:00</b>	<b>Break</b>
<b>11:00 - 11:40</b>	<b>Keynote 3</b>  <b>Title:</b> <i>Substrait: Rethinking DBMS Composability</i>  <b>Speaker:</b> <i>Jacques Nadeau, Co-creator of Apache Arrow, Dremio and Substrait</i>
<b>11:40 - 12 noon</b>	<b>Invited Talk 2</b>  <b>Title:</b> <i>GoogleSQL: A SQL Language as a Component</i>  <b>Speaker:</b> <i>Jeff Shute, Distinguished Engineer, Google</i>

<b>12 noon - 12:20</b>	<b>Invited Talk 3</b>  <b>Title:</b> <i>The Single Product Approach: Supporting Diverse Workloads with Snowflake's Query Engine</i>  <b>Speaker:</b> <i>Jiaqi Yan, Principal Software Engineer, Snowflake Computing</i>
<b>12:20 - 12:30</b>	<b>Short Paper 2</b>  <b>Title:</b> <i>Hamilton: a modular open source declarative paradigm for high level modeling of dataflows</i>  <b>Speaker:</b> <i>Stefan Krawczyk, Stitch Fix</i>
<b>12:30 - 14:00</b>	<b>Lunch</b>
<b>14:00 - 14:40</b>	<b>Keynote 4</b>  <b>Title:</b> <i>Shared, High-Performance Software Components for Shared, High-Performance Hardware</i>  <b>Speaker:</b> <i>Xiaosong Ma, Principal scientist, Qatar Computing Research Institute, HBKU</i>
<b>14:40 - 15:00</b>	<b>Invited Talk 4</b>  <b>Title:</b> <i>Coral: A SQL translation and rewrite engine for modern data lakes</i>  <b>Speaker:</b> <i>Walaa Eldin Moustafa, Senior Staff Software Engineer, LinkedIn</i>
<b>15:00 - 15:20</b>	<b>Invited Talk 5</b>  <b>Title:</b> <i>Delta Lake: The one storage format for all your analytics</i>  <b>Speaker:</b> <i>Tathagata Das, Staff Software Engineer, Databricks, Apache Spark PMC</i>
<b>15:20 - 15:30</b>	<b>Short Paper 3</b>  <b>Title:</b> <i>Efficient Property Projections of Graph Queries over Relational Data</i>  <b>Speaker:</b> <i>Danica Porobic, Oracle</i>
<b>15:30 - 16:00</b>	<b>Break</b>
<b>16:00 - 16:40</b>	<b>Keynote 5</b>  <b>Title:</b> <i>DuckDB - A Modern Modular &amp; Extensible Database System</i>  <b>Speaker:</b> <i>Mark Raasveldt, CTO, DuckDB Labs &amp; Postdoctoral Researcher, CWI</i>
<b>16:40 - 17:00</b>	<b>Invited Talk 6</b>  <b>Title:</b> <i>Velox: Unifying execution engines across Meta and beyond</i>  <b>Speaker:</b> <i>Pedro Pedreira, Software Engineer, Meta Platforms Inc.</i>

<b>17:00 - 17:10</b>	<b>Short Paper 4</b> <b>Title:</b> <i>Unifying Metadata for Stream and Batch Queries in a Cloud Service</i> <b>Speaker:</b> <i>Daniel Bauer, IBM Research Europe</i>
<b>17:10 - 17:20</b>	<b>Short Paper 5</b> <b>Title:</b> <i>Towards a Modular Data Management System Framework</i> <b>Speaker:</b> <i>Haralampos Gavriilidis, Technische Universität Berlin</i>
<b>17:20 - 17:30</b>	<b>Short Paper 6</b> <b>Title:</b> <i>Knowledge-based Recommendation for Polyglot Persistence</i> <b>Speaker:</b> <i>Philipp Eisenhuth, University of Bayreuth</i>
<b>17:30 - 17:45</b>	<b>Closing Remarks - Mohamed &amp; Satya</b>