

Schedule for First International Workshop on Composable Data Management Systems, 9th September 2022 @ VLDB

8:30 - 8:40	Welcome and Kickoff - Mohamed & Satya
8:40 - 9:20	Keynote 1 Title: <i>From parsing, to query processing, to resiliency: Deriving unexpected value from database technology through componentization</i> Speaker: <i>Jonathan Goldstein, Researcher, Microsoft</i>
9:20 - 10:00	Keynote 2 Title: <i>Evolving an Organization's Data Management Discipline in support of their Machine Learning Activities</i> Speaker: <i>David Cohen, Senior Principal Engineer, Intel</i>
10:00 - 10:20	Invited Talk 1 Title: <i>SAP's global data mesh for composable data management supporting cross-domain consumption</i> Speakers: <i>Anil K Goel, CTO / Head of Technology Office, SAP HANA Database and Analytics & Mihnea Andrei, SAP HANA Database and Analytics</i>
10:20 - 10:30	Short Paper 1 Title: <i>Assembling a Query Engine From Spare Parts</i> Speakers: <i>Mosha Pasumansky, Firebolt Analytics & Benjamin Wagner, Firebolt Analytics</i>
10:30 - 11:00	Break
11:00 - 11:40	Keynote 3 Title: <i>Substrait: Rethinking DBMS Composability</i> Speaker: <i>Jacques Nadeau, Co-founder & CEO Sundeck</i>
11:40 - 12 noon	Invited Talk 2 Title: <i>GoogleSQL: A SQL Language as a Component</i> Speaker: <i>Jeff Shute, Distinguished Engineer, Google</i>

12 noon - 12:20	Invited Talk 3 Title: <i>The Single Product Approach: Supporting Diverse Workloads with Snowflake's Query Engine</i> Speaker: <i>Jiaqi Yan, Principal Software Engineer, Snowflake Computing</i>
12:20 - 12:30	Short Paper 2 Title: <i>Hamilton: a modular open source declarative paradigm for high level modeling of dataflows</i> Speaker: <i>Stefan Krawczyk, Stitch Fix</i>
12:30 - 14:00	Lunch
14:00 - 14:40	Keynote 4 Title: <i>Shared, High-Performance Software Components for Shared, High-Performance Hardware</i> Speaker: <i>Xiaosong Ma, Principal scientist, Qatar Computing Research Institute, HBKU</i>
14:40 - 15:00	Invited Talk 4 Title: <i>Coral: A SQL translation and rewrite engine for modern data lakes</i> Speaker: <i>Walaa Eldin Moustafa, Senior Staff Software Engineer, LinkedIn</i>
15:00 - 15:20	Invited Talk 5 Title: <i>Delta Lake: The one storage format for all your analytics</i> Speaker: <i>Tathagata Das, Staff Software Engineer, Databricks, Apache Spark PMC</i>
15:20 - 15:30	Short Paper 3 Title: <i>Efficient Property Projections of Graph Queries over Relational Data</i> Speaker: <i>Danica Porobic, Oracle</i>
15:30 - 16:00	Break
16:00 - 16:40	Keynote 5 Title: <i>DuckDB - A Modern Modular & Extensible Database System</i> Speaker: <i>Mark Raasveldt, CTO, DuckDB Labs & Postdoctoral Researcher, CWI</i>
16:40 - 17:00	Invited Talk 6 Title: <i>Velox: Unifying execution engines across Meta and beyond</i> Speaker: <i>Pedro Pedreira, Software Engineer, Meta Platforms Inc.</i>

17:00 - 17:10	Short Paper 4 Title: <i>Unifying Metadata for Stream and Batch Queries in a Cloud Service</i> Speaker: <i>Daniel Bauer, IBM Research Europe</i>
17:10 - 17:20	Short Paper 5 Title: <i>Towards a Modular Data Management System Framework</i> Speaker: <i>Haralampos Gavriilidis, Technische Universität Berlin</i>
17:20 - 17:30	Short Paper 6 Title: <i>Knowledge-based Recommendation for Polyglot Persistence</i> Speaker: <i>Philipp Eisenhuth, University of Bayreuth</i>
17:30 - 17:45	Closing Remarks - Mohamed & Satya