

# Towards a Modular Data Management System Framework

**Haralampos Gavriilidis, Lennart Behme**, Sokratis Papadopoulos,  
Stefano Bortoli, Jorge-Arnulfo Quiané-Ruiz, Volker Markl

September 9, 2022



# A Brief History of Data Management Systems

- DMSes historically have a monolithic architecture
- Many special-purpose DMSes emerged since “one size does not fit all”
- The general architecture for query processors has not changed much since the 1980s though

⇒ Engineers reinvent the wheel as they have to build all components from scratch for each system

⇒ DMS components should be composable

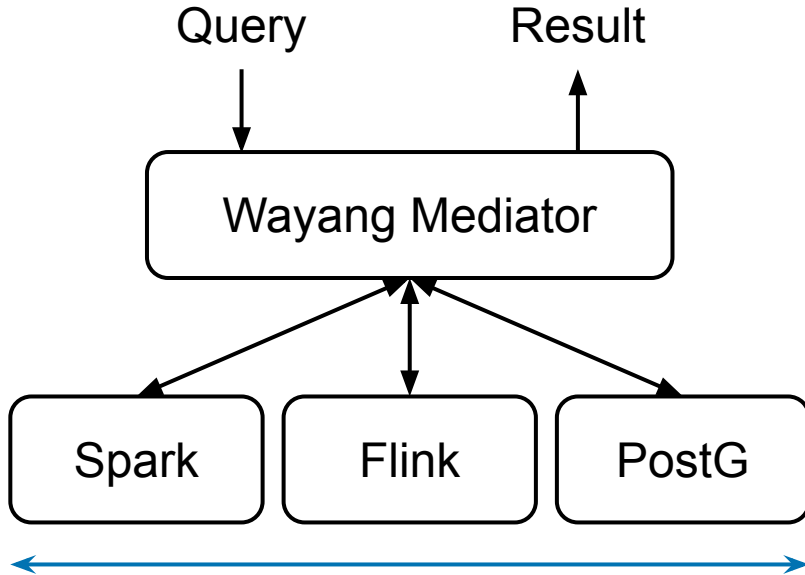
# The State of the (Composability) Union

- There are no standards that allow system engineers to make their components easily composable
- Composing existing components into a DMS requires a huge amount of manual effort
- Approaches like Presto are impractical and costly as they continuously port data and pipeline logic from legacy DMSes to the latest and most advanced system

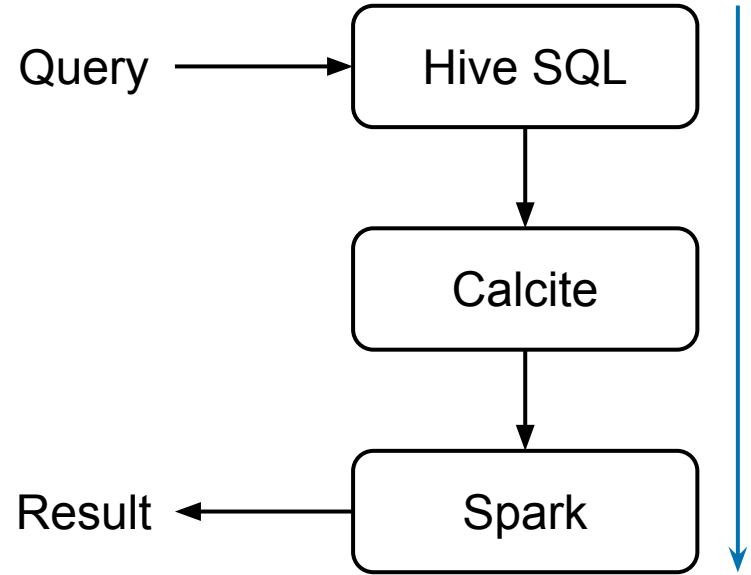
# PolyDMS – Our Vision to Address the Composability Gap

- System architects should be able to build a new DMS by reusing individual components as building blocks
- We define abstract types of DMS components, such as query interfaces or optimizers, and design generic interfaces for each type so that existing libraries can easily turn into a DMS component
- We build a Component Orchestrator that interacts with all components, enabling the composition of DMSES
- We design a high-level System Definition Language (SDL) that acts as an interface to the orchestrator and allows users to easily compose a new DMS out of existing components

# Existing DMS Composition Approaches



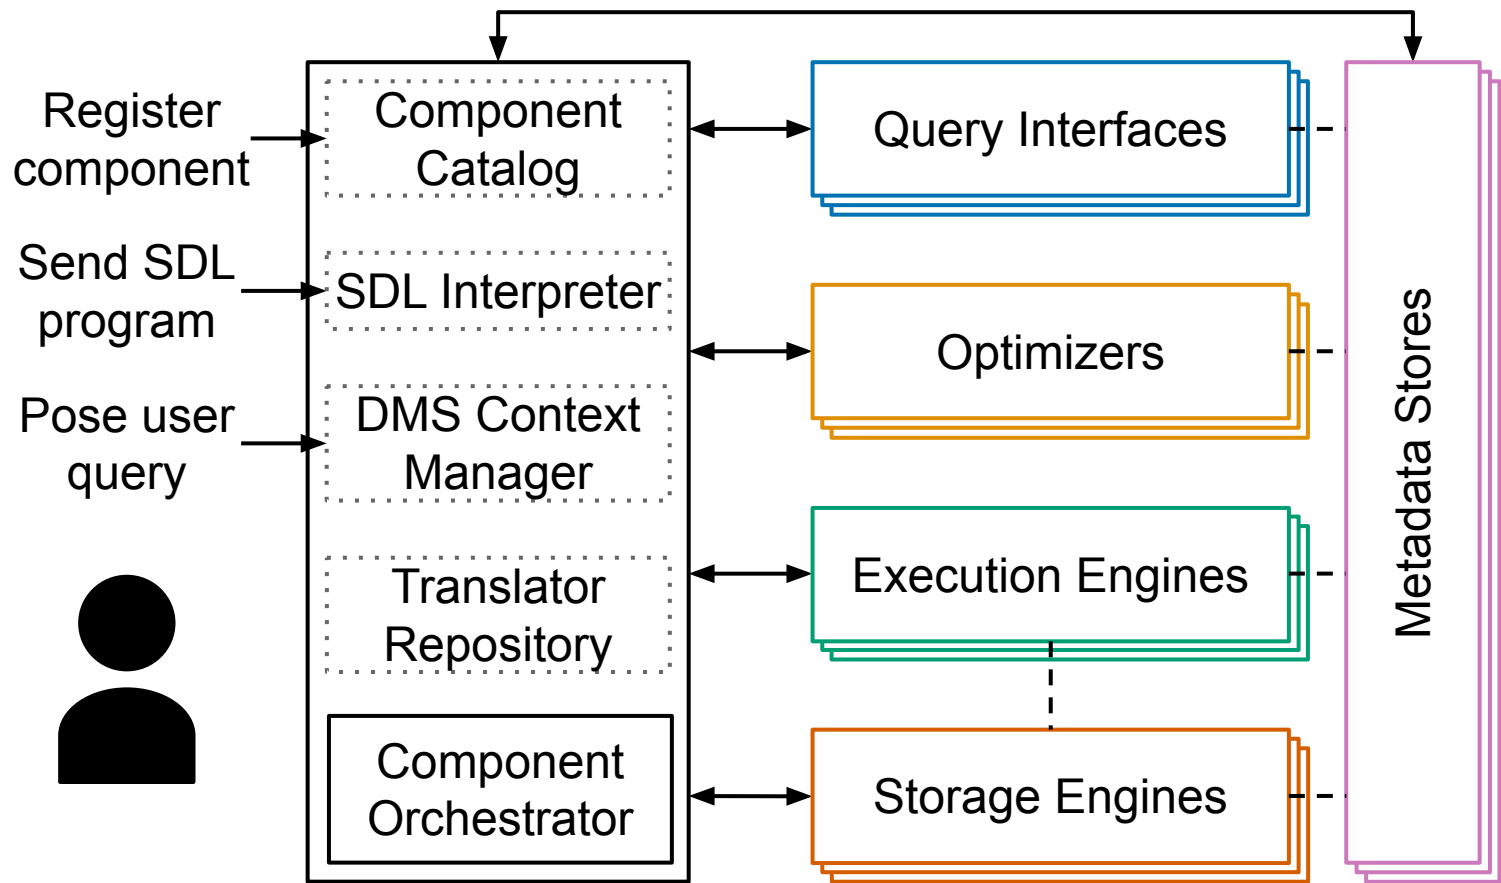
**Horizontal** (e.g., Wayang)



**Vertical** (e.g., Hive)

# PolyDMS Framework

- Existing (and future) DMS components should be independently reusable
- We propose to couple components through a micro-service architecture with well-defined APIs



# Component Types

- Metadata Stores
- Query Interfaces
- Optimizers
- Execution Engines
- Storage Engines



# Component Orchestration

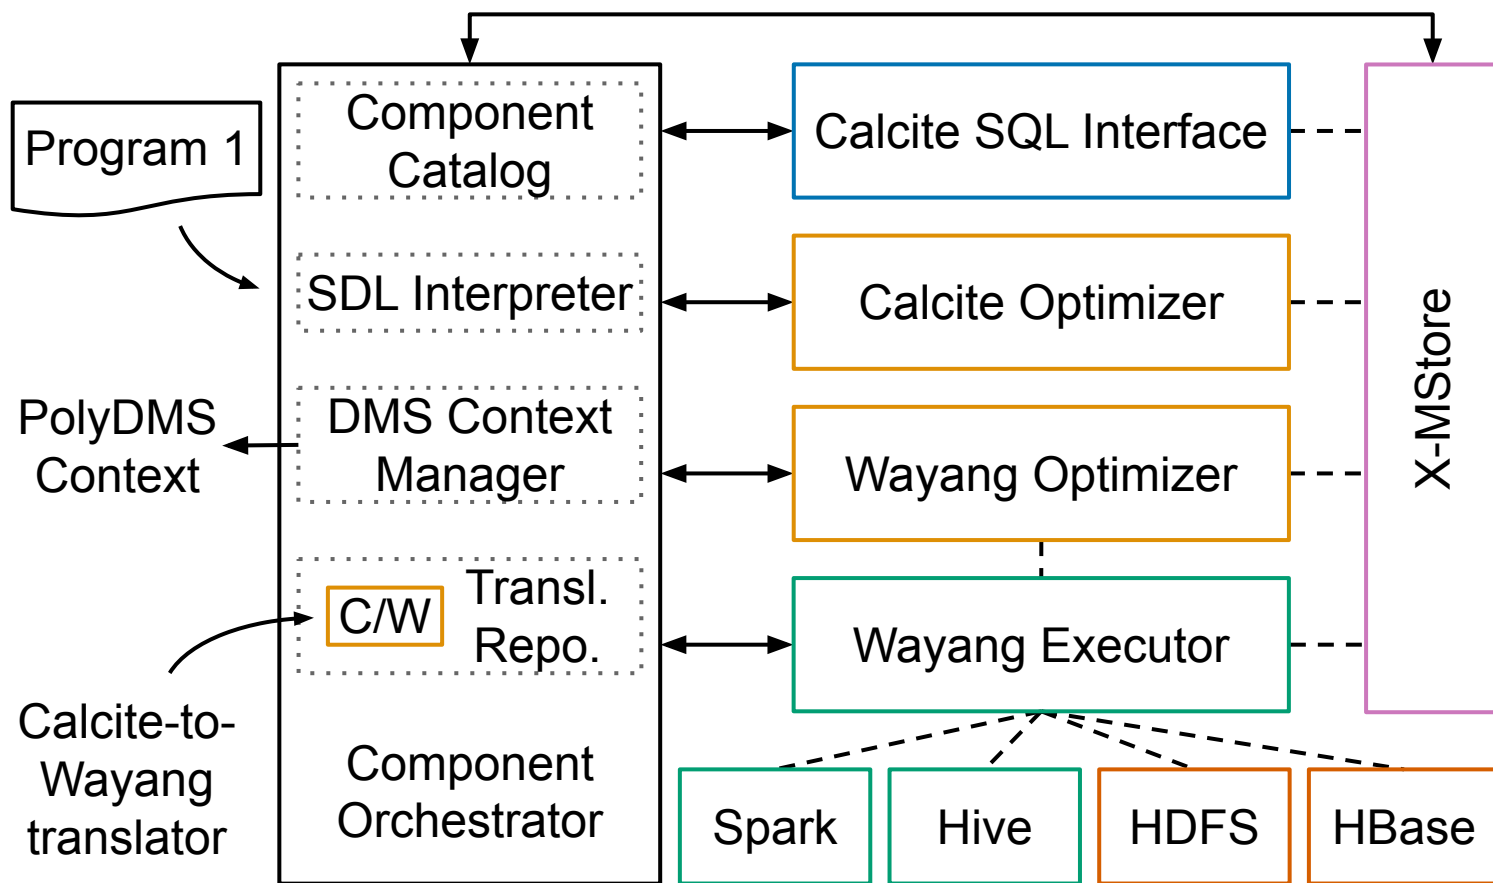
- Component Catalog
- System Definition Language (SDL) Interpreter
- DMS Context Manager
- Translator Repository

# Application Scenarios

- Diverse user requirements
- Optimization synergies

# Proof-of-Concept (PoC)

- Real-world industry use case that existing DMSes cannot handle efficiently in cooperation with Huawei
- Solution architects often face fixed legacy DMS infrastructures that require a cross-platform system for utilizing multiple heterogeneous DMSes in one workload
- Existing cross-platform systems do not support SQL interfaces or important logical optimizations like join ordering




# PoC SDL Example

---


```
def cross_platform_dms(input):  
    ctx = PolyDMSContext()  
    ctx.md = XMStore()  
    ctx.qi = CalciteSqlInterface(input, ctx.md)  
    ctx.log_opt = CalciteOpt(ctx.qi, ctx.md)  
    ctx.cp_opt = WayangOpt(translate(ctx.log_opt), ctx.md)  
    ctx.exec = WayangExec(ctx.cp_opt, ctx.md)  
  
    return ctx.initialize()
```

---

# Preliminary Evaluation



Optimization synergies



Metadata impact

# Outlook

- PolyDMS is our vision to break today's monolithic DMS architectures
- PolyDMS composes stand-alone components into custom DMSes
- Next: formalize our ideas for SDL and extend the orchestrator prototype to construct DMS instances based on SDL statements



**Haralampos Gavriilidis**

[gavriilidis@tu-berlin.de](mailto:gavriilidis@tu-berlin.de)

[www.user.tu-berlin.de/harry\\_g](http://www.user.tu-berlin.de/harry_g)



**Lennart Behme**

[lennart.behme@tu-berlin.de](mailto:lennart.behme@tu-berlin.de)

[www.lbeh.me](http://www.lbeh.me)